

Assignment 1

Fit a polynomial model.

David Padilla Orenga
Ignacio Pastore Benaim

September 30, 2024

1 Introduction

In this report, we present the results of the first assignment of the Machine Learning course. The assignment consists of fitting several regression models to a noisy dataset generated from a polynomial function. The figure 1 shows the dataset which contains 500 data points. The majority of the methods seen in class were explored for the sake of completeness. A comparison of the results is presented along with the election of the optimal models.

We explore multiple methods, including k-NN, linear regression variants, and regularized polynomial regressions. The optimal models were selected based on cross-validation, and final test results are reported. Finally, a short discussion and conclusion are presented.

The complete code used for this analysis can be found in a Google Colab Notebook.

2 Materials & Methods

The original dataset consists of 500 data points, which were split into training (81%), validation (9%), and test (10%) sets. No outlier removal was performed. The evaluation process involved 3 key steps: initial screening with validation, cross-validation of selected models, and final evaluation on the test set.

All models were implemented in Python with the numpy, pandas and scikit-learn libraries.




Figure 1: Polynomial function generated with random noise.

2.1 Initial Screening with Validation set

Various models (described in Section 2.4) were first evaluated using a validation set. The purpose of this step was to quickly assess model performance and identify promising candidates.

2.2 Cross-Validation of Selected Models

Based on the validation results, a subset of models was chosen for cross-validation with $k = 10$ folds. The average performance across the folds was used to further narrow down the best-performing models.

2.3 Final Testing on Test Set

After cross-validation, the best models were retrained with the training and validation set combined. The evaluation was finally made on the test set.

2.4 Models

We tested several different models, including k-Nearest Neighbors (k-NN), various linear regression models, and polynomial regression with regularization.

2.4.1 k-NN

The k-NN model was implemented with an optimal $k = 17$ obtained from a grid search with 5-fold cross-validation over the range of 1 to 20.

2.4.2 Linear Regression Variants

Despite knowing that a linear model might not be the outmost appropriated for this task, 3 variants were implemented to compare with other methods:

- Standard Linear Regression.
- Huber Regression.
- RANSAC Regression.

2.4.3 Polynomial Regression with Regularization

Polynomial regression models were fitted using Ridge and Lasso regularizations to control overfitting. We applied 5-fold cross-validation using a grid search over polynomial degrees from 1 to 12 and regularization strengths λ in the range of [0.01, 0.1, 1, 10, 100] to find the optimal configurations. The initial results showed that the best-performing models were Ridge Regression with degree 12 and $\lambda = 0.01$ and Lasso Regression with degree 6 and $\lambda = 0.01$.

However, visual inspection suggested that high-degree polynomial models were prone to overfitting. To balance model complexity and performance, we decided to focus on polynomial

Table 1: Initial validation performance

Model	MSE	R2
KNN (k=17)	0.514	0.916
Linear	0.978	0.841
Huber	1.059	0.828
RANSAC	1.337	0.782
Polynomial (degree=6)	0.425	0.931
Ridge (degree=6, $\lambda = 0.01$)	0.430	0.930
Lasso (degree=6, $\lambda = 0.01$)	0.474	0.923
Polynomial (degree=7)	0.509	0.917
Ridge (degree=7, $\lambda = 0.01$)	0.465	0.924
Lasso (degree=7, $\lambda = 0.01$)	0.474	0.923

degrees 6 and 7 for further analysis, both with and without regularization.

2.5 Evaluation Metrics

Through all the phases, the models were evaluated using 2 primary metrics:

- Mean Squared Error (MSE).
- R^2 Score.

3 Results

3.1 Initial Validation Results

The first screening of models was done using the validation set. Table 1 shows the MSE and R^2 scores for the various models during the validation phase. Based on these results, linear models were discarded for further analysis.

3.2 Cross-Validation Results

The cross-validation results are summarized in Table 2 showing the MSE and R^2 scores. Based on these results all the models were retained for further analysis.

Table 2: Cross-validation performance

Model	MSE	R2
Polynomial (degree=6)	1.770	0.682
Ridge (degree=6, $\lambda = 0.01$)	1.769	0.682
Lasso (degree=6, $\lambda = 0.01$)	1.813	0.676
Polynomial (degree=7)	1.763	0.684
Ridge (degree=7, $\lambda = 0.01$)	1.765	0.683
Lasso (degree=7, $\lambda = 0.01$)	1.813	0.676
KNN (k=17)	1.803	0.705

Table 3: Tests performance

Model	MSE	R2
KNN (k=17)	1.127	0.807
Polynomial (degree=6)	1.160	0.801
Ridge (degree=6, $\lambda = 0.01$)	1.155	0.802
Lasso (degree=6, $\lambda = 0.01$)	1.153	0.802
Polynomial (degree=7)	1.160	0.801
Ridge (degree=7, $\lambda = 0.01$)	1.155	0.802
Lasso (degree=7, $\lambda = 0.01$)	1.153	0.802

3.3 Test Set Results

Finally, the selected models were tested on the test set. Table 3 shows the final MSE and R^2 scores on the test set.

4 Discussion

The initial screening confirmed that linear models were unsuitable and were discarded as expected. Surprisingly, k-NN performed similarly to the more complex polynomial models. The MSE values ranged from 1.127 to 1.160, and when rounding up the R^2 scores, all models explained around 80% of the variance in the data.

Ridge and Lasso regularizations showed no significant improvements, suggesting that overfitting was not an issue. Given the comparable results, we would choose between k-NN and polynomial regression for their simplicity. Future work could involve removing outliers to refine the results further.

5 Conclusion

All tested models—k-NN, polynomial regression, and regularized variants—showed comparable performance, explaining approximately 80% of the variance in the data. Minor differences in the metrics indicate no clear best model, making k-NN and polynomial regression preferable for their simplicity. Future work could focus on enhancing model performance through data preprocessing techniques such as outlier removal.